

CYENG 312/GECE 594 :

Trusted Operating System (OS)

Lecture 03: SELinux and AppArmor

Instructor: Shayan (Sean) Taheri, Ph.D.

Assistant Professor

The Department of Electrical and Cyber Engineering (ECE)

The Institute for Health and Cyber Knowledge (I-HACK)

The Gannon University (GU)





Personal Information

- ❑ Name: Shayan (Sean) Taheri.
- ❑ Date of Birth: July/28/1991.
- ❑ Past Position: Postdoctoral Fellow at University of Florida.
- ❑ Ph.D. Degree: Electrical Engineering from the University of Central Florida.
- ❑ M.S. Degree: Computer Engineering from the Utah State University.
- ❑ University Profile:
<https://www.gannon.edu/FacultyProfiles.aspx?profile=taheri001>



SELinux and AppArmor

- Two major Open Source Secure Oses: More Security Vs. More Usability
- SELinux
 - ❑ *Security-Enhanced Linux (SELinux) is a Linux feature that provides a variety of security policies for Linux kernel.*
 - ❑ *It is included with CentOS / RHEL / Fedora Linux, Debian / Ubuntu, Suse, Slackware and many other distributions.*
 - ❑ *For Strict security, but hard to use.*
 - ❑ *Developed by NSA.*
- AppArmor
 - ❑ *AppArmor (Application Armor) is another security software for Linux which maintained and released by Novell under GPL.*
 - ❑ *AppArmor was created as an alternative to SELinux.*
 - ❑ *AppArmor works with file paths.*
 - ❑ *AppArmor is default in OpenSUSE and Suse Enterprise Linux. It was first successfully packaged for Ubuntu Linux.*
 - ❑ *Easy to use, but to not strict in security.*
 - ❑ *Was called Subdomain, developed by Immunix.*
 - ❑ *Now maintained by Novell.*



SELinux and AppArmor (Contd.)

➤ SELinux Features

- ❑ *Clean separation of policy from enforcement*
- ❑ *Well-defined policy interfaces*
- ❑ *Support for applications querying the policy and enforcing access control*
- ❑ *Independent of specific policies and policy languages*
- ❑ *Independent of specific security label formats and contents*
- ❑ *Individual labels and controls for kernel objects and services*
- ❑ *Caching of access decisions for efficiency*
- ❑ *Support for policy changes*
- ❑ *Separate measures for protecting system integrity (domain-type) and data confidentiality (multilevel security)*
- ❑ *Very flexible policy*
- ❑ *Controls over process initialization and inheritance and program execution*
- ❑ *Controls over file systems, directories, files, and open file descriptors*
- ❑ *Controls over sockets, messages, and network interfaces*
- ❑ *Controls over use of “capabilities”*



SELinux and AppArmor (Contd.)

➤ AppArmor Features

- ❑ *Full integration.*
- ❑ *Easy deployment.*
- ❑ *AppArmor includes a full suite of console and YaST-based tools to help you develop, deploy and maintain application security policies.*
- ❑ *Protects the operating system, custom and third-party applications from both external and internal threats by enforcing appropriate application behavior.*
- ❑ *Reporting and alerting. Built-in features allow you to schedule detailed event reports and configure alerts based on user-defined events.*
- ❑ *Sub-process confinement. AppArmor allows you to define security policies for individual Perl and PHP scripts for tighter Web-server security.*



SELinux and AppArmor (Contd.)

➤ SELinux Pros and Cons

- ❑ *Admin skill set (learning curve) – High*
- ❑ *Complex and powerful access control mechanism – Yes*
- ❑ *Detailed configuration required – Yes*
- ❑ *GUI tools to write / modify rules set – Yes*
- ❑ *CLI tools to write / modify rules set – Yes*
- ❑ *Ease of use – No (often described as horrible to use)*
- ❑ *Binary package – Available for most Linux distributions*
- ❑ *System performance impact: None*
- ❑ *Security Framework: Mandatory access controls using Flask*
- ❑ *Auditing and logging supported – Yes*
- ❑ *Typical user base – Enterprise users*
- ❑ *Documentation – Well documented*



SELinux and AppArmor (Contd.)

➤ AppArmor Pros and Cons

- ☐ Admin skill set (learning curve) – Medium
- ☐ Complex and powerful access control mechanism – Yes.
- ☐ Detailed configuration required – Yes.
- ☐ GUI tools to write / modify rules set – Yes (yast2 and wizards).
- ☐ CLI tools to write / modify rules set – Yes.
- ☐ Ease of use – Yes (often described as less complex and easier for the average user to learn than SELinux).
- ☐ Binary package – Available for Ubuntu / Suse / Opensuse and distros.
- ☐ System performance impact – None.
- ☐ Security Framework – Mandatory access controls.
- ☐ Auditing and logging supported – Yes.
- ☐ Typical user base – Enterprise users.
- ☐ Documentation – Documented (mostly available from Opensuse and Suse enterprise Linux).



Access Control in SELinux: Type Enforcement

- Label based access control
 - **Domain** Identifier for process.
 - **Type** Identifier (label) for resources.
 - Controls permission between domain and type.
- Fine-grained access control
 - A method of controlling who can access certain data.
 - Compared to generalized data access control, also known as coarse-grained access control, fine-grained access control uses more nuanced and variable methods for allowing access.





SELinux: Configuration of policy

- The most important feature
 - What domain can access what access to what types?

Ex. Web server: domain httpd_t: Allowing access to homepage

◆ **allow specify domain, type, permission**

◆ **allow** httpd_t web_contents_t file:{ read };

Domain

Type

Permission

◆ Assign label (= type) to resource

/var/www(|/.*) **system_u: object_r:web_contents_t**

- Many lines of allows (10k-100k) are required
- Macro is Used: Bunch of allows is summarized by macro



Reminder - Macro

- It is a rule or pattern that specifies how a certain input should be mapped to a replacement output.
- Applying a macro to an input is known as macro expansion.
- The input and output may be a sequence of lexical tokens or characters, or a syntax tree.
- Macros are used to make a sequence of computing instructions available to the programmer as a single program statement, making the programming task less tedious and less error-prone.



Example of policy

●bind.te: allowing acces

```

type named_t;
type named_exec_t;
init_daemon_domain(named_t,named_exec_t)

...
...
kernel_read_kernel_sysctls(named_t)
kernel_read_system_state(named_t)
kernel_read_network_state(named_t)
kernel_tcp_recvfrom(named_t)
....
corenet_tcp_sendrecv_all_if(named_t)
corenet_raw_sendrecv_all_if(named_t)
corenet_udp_sendrecv_all_if(named_t)
corenet_tcp_sendrecv_all_nodes(named_t)
corenet_udp_sendrecv_all_nodes(named_t)
corenet_raw_sendrecv_all_nodes(named_t)
corenet_tcp_sendrecv_all_ports(named_t)
corenet_udp_sendrecv_all_ports(named_t)
corenet_non_ipsec_sendrecv(named_t)
corenet_tcp_bind_all_nodes(named_t)
corenet_udp_bind_all_nodes(named_t)

```

●bind.fc:assigning label

```

/etc/rndc.* -- gen_context(system_u:object_r:named_conf_t,s0)
/etc/rndc.key -- gen_context(system_u:object_r:dnsssec_t,s0)

/usr/sbin/lwresd -- gen_context(system_u:object_r:named_exec_t,s0)
/usr/sbin/named -- gen_context(system_u:object_r:named_exec_t,s0)
/usr/sbin/named-checkconf -- gen_context(system_u:object_r:named_checkconf_exec_t,s0)
/usr/sbin/r?ndc -- gen_context(system_u:object_r:ndc_exec_t,s0)

/var/log/named.* -- gen_context(system_u:object_r:named_log_t,s0)

/var/run/ndc -s gen_context(system_u:object_r:named_var_run_t,s0)
/var/run/bind(/.*)? gen_context(system_u:object_r:named_var_run_t,s0)
/var/run/named(/.*)? gen_context(system_u:object_r:named_var_run_t,s0)

ifdef(`distro_debian',`
/etc/bind(/.*)? gen_context(system_u:object_r:named_zone_t,s0)
/etc/bind/named.conf -- gen_context(system_u:object_r:named_conf_t,s0)

```

**...45
labels**

...293 lines

...100 kinds of macros

Difficult to understand



Access Control in AppArmor

- Easier than SELinux Implemented as Loadable Kernel Module (LKM).
- Recently, often compared with SELinux.
- Features:
 - Access control
 - ✓ Controls file and POSIX capability
 - ✓ Path name-based
 - Label is not used
 - ✓ Profile → “policy”
 - GUI Tools
 - ✓ Integrated in YaST
 - Generating profile
 - Log report
 - Not so important for embedded computing



AppArmor Path Name-Based Access Control

➤ Path name based:

- Identify file with “path name”
- Easy to understand

➤ Example:

```
/usr/sbin/httpd{
```

```
  /var/www/** r,
```

```
}
```

→ /usr/sbin/httpd **can read under** /var/www



Permission to File

- Basic permission: r,w,x,l
 - r read
 - w : write
 - ix : execute
 - l : link (remove file)



POSIX Capability

- Controls capability
 - Capability
 - ✓ Important operation other than file access
 - ✓ Example:
 - net_bind_service: bind well-known port
 - net_raw: use raw socket
 - For detail: see `$mancapabilities`
- The **Portable Operating System Interface (POSIX)** is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems.
- POSIX defines both the system- and user-level application programming interfaces (API), along with command line shells and utility interfaces, for software compatibility (portability) with variants of Unix and other operating systems.
- POSIX is intended to be used by both application and system developers.



Configuration for Profile

- Simple, easy to understand

```
/usr/sbin/named {      -> path to  
executable
```

```
#include <abstractions/base>
```

```
#include<abstractions/nameservice>
```

} Common

```
capability net_bind_service,
```

```
capability setgid,
```

```
capability setuid,
```

} Capability

```
<snip>
```

```
/var/lib/named/** rwl,
```

```
/var/run/named.pid wl,
```

} Access to file

```
}
```



Linux Security Module (LSM)

- Both use LSM for implementation
- LSM: Linux Security Module
 - Set of hooks in kernel to check security
 - It is included in mainline from LIDS 2.6 version
- Using LSM:
 - SELinux, AppArmor, LIDS (for 2.6 version)
- Not using
 - TOMOYO Linux, LIDS (for 2.4 version)



Difference Between SELinux and AppArmor

- Granularity of permission
 - SELinux:
 - ✓ File, network, IPC, POSIX, capability, etc.
 - AppArmor
 - ✓ File + POSIX capability
 - AppArmor can reach SELinux in theory, because both use LSM.



How to identify resource

- Fundamental difference
 - Affects security and usability
- Label based Vs. Path name based
 - Label: lower usability, higher security
 - Assign label to file
 - SELinux
 - Path name: higher usability, lower security
 - Identify file with path name
 - AppArmor, TOMOYO Linux
- Compare them by showing benefit and loss of pathname



Benefit of Path Name-Based

- High usability, easy to understand → No
- need to extend file system
 - Label base: File system have to be extended to store label.
- Implementing policy generation tool is easier.
- Nothing happens when index node (inode) number is changed.




Benefit of Path Name-Based in Policy Generation

- Example case:
 - PHP tried to write /var/www/html/write/test.txt → But, access denied by Secure OS
 - Have to generate policy from log
- SELinux
 - 1) label under /var/www/html -> httpd_sys_content_t
 - 2) Log says:
 - *httpd_t was denied to write to httpd_sys_content_t*
 - 3) Generate policy from log:
 - *allow httpd_t httpd_sys_content_t: file write;*
- > **allowing write access for whole “/var/www” resource!**
 - 4) Unnecessary access is granted.**
- AppArmor
 - 1) log says:
 - /usr/sbin/httpd is denied to write /var/www/html/write/test.txt
 - 2) Generate policy (= profile) from log:
 - **/usr/sbin/httpd{/var/www/html/write/test.txt w}**
 - 3) Unnecessary access is NOT granted → Benefit of Path Name-Based in AppArmor.**



Reminder – “index node (inode)”

- The inode (index node) is a data structure in a Unix-style file system that describes a file-system object such as a file or a directory.
- Each inode stores the attributes and disk block locations of the object's data.
- File-system object attributes may include metadata (times of last change, access, modification), as well as owner and permission data.
- A directory is a list of inodes with their assigned names. The list includes an entry for itself, its parent, and each of its children.



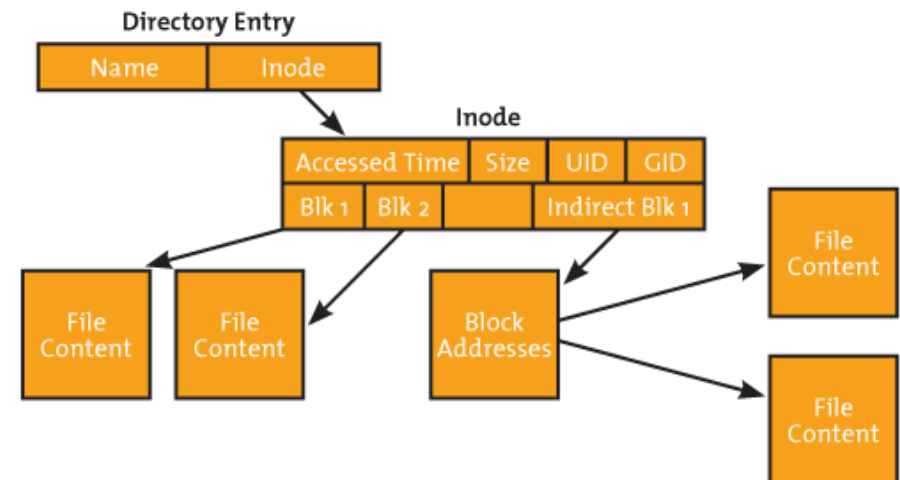
Directory /home/user

File Name	Inode Number
File1.txt	123
File2.txt	124
File9.txt	119

- File types
- Permissions
- UID
- GID
- FileSize
- Data block

Blocks/disks

Data
Data
Data





Benefit of Path Name Based in Change of inode Number

- Example: “/etc/mtab”.
- SELinux: Label is lost when inode number is changed
 - Label is associated with inode
 - /etc/mtab
 - **vi** (Visual Instrument, a text editor), **rpm** (Red Hat Package Manager, a command-line utility for managing packages) changes inode
 - Solution
 - “file type transition” configuration
 - Not easy for beginner
 - Some userland have to be extended, Example: rpm ,vi.
- AppArmor:
 - No problem!



Loss by Path Name-Based

- **Loss By:**
 - Information Flow Analysis
 - Temporary (tmp) Files
 - Who can access the information?
- Some people say path name-based security is broken because of:
- **Ex:** Information flow analysis to password information.
 - Initial State: Stored in → /etc/shadow
 - If hardlink is created to /etc/shadow, password information can be accessed via hardlink.
 - In computing, a hard link is a directory entry (in a directory-based file system) that associates a name with a file.
 - Thus, each file must have at least one hard link.
 - What happens in information flow analysis?
 - Have to traverse whole file tree to find hardlink.
 - What if more hardlink is created during traversal?
 - SELinux:
 - All you have to do is to check what kind of domain can access label for /etc/shadow.
 - Label is the same for hardlink.



Loss by Path Name-Based in Temporary Files

- When creating randomly named file under /tmp
- SELinux
 - Can identify such file by naming label such as httpd_tmp_t.
- AppArmor
 - How to identify randomly named files?
 - ✓ Result in allowing whole /tmp.



SELinux Policy Editor (SEEDIT)

- Tool that makes SELinux easy
- Open Source: <http://seedit.sourceforge.net/>
 - Originally developed by Hitachi Software
 - Included in Fedora repository

Main Feature: Simplified Policy Description Language (SPDL)

- AppArmor-like syntax to write policy
 - Example:
 - domain httpd_t
 - program /usr/sbin/httpd;
 - allow /var/www/** r; → **path name configuration**
 - This is converted to SELinux policy syntax
 - type var_www_t; → **label is generated**
 - allow httpd_t var_www_t { file dir }: read;



SELinux Policy Editor (Contd.)

- Still different from AppArmor
- Inherit drawback from label-based access control
 - Change of inode
 - Generated policy is label-based
- Inherit good points from SELinux
 - Fine-grained permission (interprocess communication - IPC, network)
 - IPC refers specifically to the mechanisms an operating system provides to allow the processes to manage shared data.
 - Typically, applications can use IPC, categorized as clients and servers, where the client requests data and the server responds to client requests.
 - No patch to kernel
- **Experiment**: Porting SELinux/AppArmor to embedded devices to learn more.



Questions?